

Task 1 and Task 2: Complete Evidence Report

Verified growth-machine run, SuperDocs proposal experiment, screenshots, findings, and remaining work

Submission evidence compilation

15 August 2026

Outcome. Task 1 ran successfully through the repaired n8n control plane. The repository records eight inputs, eight draft outputs, zero emails sent, zero final claim-check failures, and zero risk-flagged rows. Task 2 then tested a proposal-review workflow in SuperDocs. The first broad request failed to return the requested ranking and later produced an inaccurate change summary. A narrower retry succeeded, identified Timeline as the weakest section, and led to two reviewed Timeline changes.

Evidence area	Included in this edition
Task 1 result	Verified n8n metrics, document-pain distribution, and two new successful workflow screenshots.
Task 2 artifact	Detailed synthetic proposal, protected clauses, and its connection to the Task 1 insight.
SuperDocs test	Failure trail, unchanged-text verification, successful ranking, targeted revision, and acceptance review.
Screenshot record	Eight current screenshots embedded beside the findings they support. No obsolete n8n template screenshot is included.
Readiness	Supported conclusions, limitations, public-post evidence map, and final actions.

The report separates repository-verified facts from screenshot observations. It does not convert synthetic inputs into customer claims or infer results that the available evidence cannot prove.

Contents

1	Purpose and evidence standard	3
2	Task 1: growth machine and n8n execution	3
2.1	System purpose	3
2.2	Successful n8n result	3
2.3	Verified run metrics	4
2.4	What Task 1 found	5
3	Task 2: proposal artifact and experiment design	5
3.1	Controlled proposal	5
3.2	Protected content	6
3.3	Evaluation goal	6
4	First interaction: failure to deliver the requested ranking	6
5	Second failure: a change summary without actual changes	7
6	Prompt correction and successful ranking	10
7	Targeted Timeline revision and human review	12
8	Consolidated results	14
9	What the evidence supports	14
9.1	Supported conclusions	14
9.2	Claims that remain out of bounds	14
10	Five-post evidence sequence	15
11	Remaining work	15
12	Recommended completion statement	15
A	Prompts preserved for reproduction	15
A.1	Constrained ranking	15
A.2	Targeted Timeline revision	16
A.3	Final correction	16
B	Evidence inventory	16

1 Purpose and evidence standard

This report consolidates the work completed for Task 1 and the proposal experiment performed for Task 2. It replaces the earlier report in full. The narrative follows the actual sequence of work, and the screenshots appear in the sections where they matter.

Three labels define the strength of each claim:

Evidence type		Label	Meaning
Repository fact	arti-	VERIFIED	A metric, file, output folder, or normalized text comparison was inspected directly.
Screenshot evidence	evi-	OBSERVED	The interface visibly supports the stated event, but the screenshot alone does not prove hidden state or whole-file integrity.
Missing item	closure	PENDING	A final export, complete post-edit text, or publication URL has not yet been supplied.

The screenshots document user-visible behavior. Quantitative claims come from stored files, not from visual inference.

2 Task 1: growth machine and n8n execution

2.1 System purpose

Task 1 is a draft-only growth machine for seed-stage technical founders and founding engineers at small B2B product companies. Human research supplies structured CSV rows. The Python system validates each row, maps a document pain to an angle, creates a claim-checked Markdown draft, and records run evidence. Sending remains hard-disabled.

The n8n workflow is a control plane around that tested Python implementation. It does not duplicate or replace the source logic. Its five operational stages are manual trigger, configuration, command validation, Python execution, and a final safety gate.

2.2 Successful n8n result

The new screenshots show the repaired workflow after execution. Each operational node is green, and the connectors show one item moving through the chain. Figures 1 and 2 replace the obsolete pre-fix template image used in the earlier report.




Figure 1: New Task 1 success screenshot. All five operational nodes show green completion states, and one item reaches the Safety Gate and Run Summary node.



Figure 2: Second supplied view of the successful n8n workflow. The screenshot confirms a completed visual path; stored metrics provide the numerical result.

2.3 Verified run metrics

The stored `n8n-run-1` evidence contains a batch snapshot, metrics, per-row results, and a human-readable summary. Eight matching Markdown drafts appear in its corresponding outbox folder.

Metric	Result	Meaning
Batch size	8	Eight synthetic prospect rows entered the machine.
Drafts written	8	Every accepted row produced a draft; completion was 100%.
Emails sent	0	The no-send boundary held.
Claim-check failures	0	No final output failed the product-claim checker.
Rows with risk flags	0	No accepted row retained a recorded risk flag.
Generator mode	8 deterministic	The run used reproducible, non-LLM generation.

VERIFIED The visible successful workflow is corroborated by `metrics.json`, `results.json`, `SUMMARY.md`, and eight output drafts. The record states `send_mode=false` and `emails_sent=0`.

2.4 What Task 1 found

Customer proposals were the largest pain category in the n8n run: four of eight rows. Across both canonical batches, proposals accounted for seven of sixteen rows.

Document pain	n8n run count	Share
Customer proposal	4	50.0%
Security questionnaire	1	12.5%
Investor update	1	12.5%
Policy pack	1	12.5%
Launch one-pager	1	12.5%

This distribution justified using a proposal as the Task 2 test artifact. It does not establish market demand, response rates, or customer outcomes. No outreach occurred.

3 Task 2: proposal artifact and experiment design

3.1 Controlled proposal

The test document is a synthetic Workflow Modernization Proposal prepared for Atlas Field Systems by Northstar Workflow Studio. Its 1,910 words cover the executive summary, objectives, current state, three-phase approach, six deliverables, roles, working principles, six-week Timeline, acceptance criteria, assumptions, scope limits, risks, change control, pricing, terms, and next steps.

The source files are:

- `atlas-field-systems-proposal.md`;
- `atlas-field-systems-detailed-proposal.docx`.

The scenario is fictional. Task 1 selected the problem category; it did not discover Atlas Field Systems or produce this proposal as a customer result.

3.2 Protected content

Two clauses created a precise integrity test.

Pricing: Pricing remains unchanged from the approved template.

Terms: This proposal does not modify the legal terms in the approved template. Any change to scope, timeline, pricing, or legal terms requires written approval from both parties.

The prompt also prohibited invented facts about systems, staffing, reporting volume, compliance, financial impact, and performance.

3.3 Evaluation goal

The test asked SuperDocs to rank sections against one approval goal: make the proposal concise, credible, and easy for an operations buyer to approve. The intended flow was analysis first, one section-level revision second, human review third, and export last.

4 First interaction: failure to deliver the requested ranking

The initial prompt requested a weakest-to-strongest ranking of every section, grounded in the document, with no edits. SuperDocs first returned a safety/rephrase response. A repeat produced several progress updates but no ranking, as Figure 3 shows.




Figure 3: Initial ranking failure. The first response asks for a rephrase; the retry supplies progress statements but still does not answer the ranking request.

OBSERVED The broad request failed at instruction completion. It did not return the requested ordered analysis.

5 Second failure: a change summary without actual changes

After the stalled ranking, the user requested a breakdown of changes. SuperDocs acknowledged that no before/after diff was available, then generated a long summary describing alleged revisions across the proposal. Figure 4 shows the beginning; Figure 5 contains the continuation.




Figure 4: Beginning of the ungrounded change summary. The response first states that no before/after diff is available, then describes sections as updated or tightened.




Figure 5: Continuation of the summary, including claims about Timeline, protected Pricing and Terms, and other sections.

The full document text was then pasted back and compared with the original source. Markdown heading markers, bullet markers, bold markers, and whitespace were normalized before comparison.

Check	SuperDocs text	Original source
Normalized words	1,910	1,910
Normalized full text	Identical	
Pricing clause	Exact match	
Terms clause	Exact match	

VERIFIED At this checkpoint, the document had not changed. The response described revisions that were absent from the actual text.

6 Prompt correction and successful ranking

The retry reduced the scope to exactly five named sections. It prohibited editing and summarizing, demanded a numbered order, required one grounded reason per section, and asked the system to name only the weakest section at the end.

This narrower contract succeeded. Figure 6 records the complete response.




Figure 6: Successful constrained ranking. SuperDocs ranked all five requested sections and named Timeline as weakest.

Rank	Section	Returned assessment
1	Timeline	Weakest; exact dates and milestones remained subject to discovery.
2	Executive Summary	Clear but high-level and cautious.
3	Next Steps	Practical but generic.
4	Proposed Approach	Structured and credible through its three phases.
5	Current State	Strongest because it stays within confirmed facts.

The result followed the requested format. Its Timeline reasoning remained imperfect because the

source already contained weekly milestones, even though calendar dates were intentionally unresolved. The output was useful, not infallible.

7 Targeted Timeline revision and human review

SuperDocs was asked to rewrite only Timeline. The instruction preserved the six-week duration and existing weekly activities, requested purpose/output/approval framing, prohibited invented facts, protected Pricing and Terms, and required review before application.

The proposed wording improved decision structure. A review caught one weakened control sentence before acceptance, so the final instruction restored the original exact wording: “The six-week duration will not change without written agreement from both parties.”

Figure 7 shows the in-document tracked revision. Removed wording appears red; additions appear green.




Figure 7: Tracked Timeline revision in the document. The first weekly heading remains visible below the edited introduction.

Figure 8 shows the review panel. Both visible edits are marked Accepted: the heading changes from Timeline to Project Timeline, and the introductory paragraph is rewritten.




Figure 8: Review panel showing two changes reviewed and both accepted. The screenshot preserves the before/after text for each visible edit.

OBSERVED The interface shows two accepted Timeline edits. The screenshot does not expose the entire post-edit file, so final whole-document integrity still requires an exported file or complete final text.

8 Consolidated results

Test	Result	Evidence
Task 1 n8n execution	Passed	New green-node screenshots plus repository metrics and eight draft files.
Draft completion	8/8	Stored <code>n8n-run-1</code> metrics.
Send safety	Passed	Zero emails sent; send mode false.
Claim safety	Passed	Zero final claim-check failures.
Broad ranking instruction	Failed	Rephrase response and unresolved progress trail.
Change-state grounding	Failed	Summary claimed changes; normalized text remained identical.
Constrained ranking	Passed	Five ordered sections with reasons and one named weakest section.
Pre-edit integrity	Passed	1,910 normalized words on both sides; full text and protected clauses matched.
Targeted review flow	Passed visually	Two visible Timeline edits reviewed and accepted.
Final exported integrity	Pending	No complete final export has been inspected.

One of two ranking formulations returned the requested output. That observation describes this small experiment only. It is not a product reliability rate.

9 What the evidence supports

9.1 Supported conclusions

- The repaired n8n workflow completed the full Task 1 control path.
- Stored evidence confirms eight drafts, zero sends, zero claim-check failures, and zero risk-flagged rows.
- Customer proposals were the largest document-pain category in the available Task 1 batches.
- The first SuperDocs interaction exposed a genuine instruction-following failure.
- The later change summary was inconsistent with the unchanged document text.
- A constrained prompt produced a complete section ranking.
- SuperDocs exposed targeted Timeline edits through a human review surface.
- The user accepted both visible edits after preserving the written-agreement control.

9.2 Claims that remain out of bounds

- The evidence does not establish customer adoption, outreach response, revenue, or time saved.
- The Atlas Field Systems scenario is synthetic, not a real account or Task 1 prospect.
- Screenshots cannot prove that every hidden or off-screen section remained unchanged after acceptance.
- Export formatting and complete post-edit integrity remain unverified.
- Two prompt attempts cannot support a general reliability percentage.

10 Five-post evidence sequence

Post	Story function	Evidence to use
1	Plan and disclosure	Explain the section-critique test, synthetic proposal, success criteria, and candidate affiliation.
2	First breakage	Use Figure 3; describe the incomplete ranking response.
3	Decision changed	Explain why the prompt was narrowed to five sections with an analysis-only output contract.
4	Measured result	Report 1,910-word pre-edit integrity, one failed broad attempt, and one successful constrained attempt.
5	Finished build	Use Figures 6, 7, and 8; add the verified final export when available.

Publish the posts across the required cadence. Record dates and URLs only after they exist.

11 Remaining work

1. Export the accepted proposal from SuperDocs as DOCX or PDF.
2. Re-open it and inspect heading hierarchy, lists, pagination, and the full Timeline.
3. Compare final Pricing and Terms with the protected strings in this report.
4. Compare all non-Timeline sections with the original source.
5. Save the final export and comparison result beside the build evidence.
6. Publish the five-post series and record genuine URLs.
7. Include a short Task 2 segment in the Task 4 demo; no separate Task 2 video is required.

12 Recommended completion statement

Task 1 ran successfully through n8n and produced eight draft-only messages with zero sends and zero final claim-check failures. Task 2 used the strongest Task 1 pain category, customer proposals, to test section critique and targeted revision in SuperDocs. The first broad ranking request failed, and a later summary described changes that had not occurred. A constrained five-section retry succeeded. SuperDocs then proposed a targeted Timeline update, and two visible changes were reviewed and accepted. Final whole-file export verification remains open.

A Prompts preserved for reproduction

A.1 Constrained ranking

Do not edit, rewrite, summarize, or propose changes to the document. Perform analysis only. Review exactly these five sections: Executive Summary, Current State, Proposed Approach, Timeline, and Next Steps. Rank them from 1, weakest, to 5, strongest against this goal: concise, credible, and easy for an operations buyer to approve. For each section, provide one reason grounded in its current wording. End by naming only the weakest section. Do not modify the document.

A.2 Targeted Timeline revision

Rewrite only the Timeline section so it is more decision-ready for an operations buyer. Preserve the six-week duration and the existing week-by-week activities. Clarify the purpose, output, and approval point for each stage. Do not invent calendar dates, costs, systems, performance targets, or customer facts. Do not modify any other section. Pricing and Terms must remain exactly unchanged. Present the proposed Timeline revision for review before applying it.

A.3 Final correction

Apply this proposed Timeline revision only, but replace the final sentence with the original exact wording: “The six-week duration will not change without written agreement from both parties.” Do not change any other section. Keep Pricing and Terms exactly unchanged.

B Evidence inventory

Evidence	Location or status
n8n metrics	task-1-growth-machine/runs/n8n-run-1/metrics.json
n8n summary	task-1-growth-machine/runs/n8n-run-1/SUMMARY.md
n8n drafts	task-1-growth-machine/outbox/n8n-run-1/
Proposal source	atlas-field-systems-proposal.md
Proposal DOCX	atlas-field-systems-detailed-proposal.docx
New Task 1 screenshots	Figures 1 and 2
SuperDocs screenshots	Figures 3–8
Embedded screenshots	Eight current images; obsolete n8n screenshots excluded
Final exported proposal	PENDING
Published post URLs	PENDING